


Task2View pipeline v2

red = new or changed from v1



WHY THESE CHANGES

Phase 3 rewritten

v1 keyword scoping, measured on GR10:
42% recall · 0 of 8 views covered
always missed the persistence boundary and the facade — named for their role, not for the task
+1 hop → **95%** · +content → **100%**

so: seed by centrality too, expand, and never cut except for budget

Phase 4 split in three

v1 could not tell a misread repo from a mis-rendered diagram — and it made CIAO vs ArchView uninterpretable

the view model fixes both, and carries evidence per element

Semantic gate added

v1 checked syntax only — ArchView already does that

the gate discriminates: a random class pair passes just **6.2%** of the time, human diagrams score **129/129**

Confidence is measured now

no invented floats — agreement is the fraction of the N parallel samples that produced the element